CSL302– Artificial Intelligence
Lab 2

Due on 17/2/2016 11.55pm

Instructions: Upload to your moodle account one zip file containing the following. Please do not submit hardcopy of your solutions. In case moodle is not accessible email the zip file to the instructor at ckn@iitrpr.ac.in. Late submission is not allowed without prior approval of the instructor. You are expected to follow the honor code of the course while doing this homework.

1. **You can work in teams of size at most 2 for this programming lab**.

2. A neatly formatted PDF document with your answers for each of the questions in the homework. You can use latex, MS word or any other software to create the PDF.

3. Include a separate folder named as 'code' containing the scripts for the homework along with the necessary data files.

4. Include a README file explaining how to execute the scripts.

5. Name the ZIP file using the following convention rollnumber1rollnumber2hwnumber.zip

1. Coal Blocks Auction (20 points)

In this problem you will be experimenting with different local search algorithms. You will take a complex problem, learn to formulate it as an AI single agent search problem and solve it using local search algorithms.

The Ministry for Power, Coal, New and Renewable Energy needs help with the auctioning of the coal blocks whose allotments were cancelled by the Supreme Court. Being the smartest AI student, the Ministry has approached you for writing software that will take as inputs the different bids and outputs an allotment that will maximize the government's revenue. The Ministry has decided to auction N coal blocks and has asked companies to bid for each block. However, the companies with aim of maximizing their profit submit bids that combine multiple blocks. The companies would like to have ownership of as many blocks in a state so as to minimize their costs for developing the infrastructure. For example, a company might bid 400 (in crores) for blocks 1, 3, 4, and 7, but bid 100 for only 1, 3 and 4.

Your software should tell the Ministry the bids that should be accepted. The Ministry has decided not to accept more than one bid from any company and not every company needs to win a bid. A company can submit any number of bids. It might also happen that some of the coal blocks do not get allotted at the end of the auction.

Formulate this problem as a state space search for a single agent and implement hill climbing with random restarts as your initial solution. You are welcome to add any novel element to the hill-climbing search. However, you are not allowed to use any integer linear programming or related algorithms to solve this problem. Include in your submission a brief description of your modification to hill climbing search.

Input Format:
Your code must be able to read inputs from a text file that will be provided as command argument. The input format is as follows

T
N
B
C
cid ncid
cid nbid bid-value block-id block id … block id
cid nbid bid-value block-id block-id … block-id

T refers to the wallclock time in minutes that will be given to your software to find a solution. *N* is the number of coal blocks; *C* is the number of companies; *B* is the number of bids. The bids of a single company are grouped together. The first line is the company id (*cid*) and the total number of bids made this company (*ncid*). Each line thereafter corresponds to a bid. Each bid starts with the company id, the number of coal blocks bid by the company in this bid (*nbid*), the bid-value, and the coal block ids. The bids are implicitly numbered. This is repeated for every company. For example,

3
7
2
6
0 3
0 4 400 0 2 3 6
0 3 100 0 2 3
0 3 50 2 3 4
1 3
1 3 300 1 4 5
1 2 200 1 5
1 3 40 1 3 4

Output Format:
The output of your code (the winning bids) should be written to a text file that will also be provide as a command argument in the following format:

revenue bid-id bid-id … bid-id

For the example problem, this would be

700 0 3

Program Execution Format:

We will be running automated scripts to check the output of your code. Hence we should be able to run your code as follows:

mysolution inputfile outputfile

Evaluation:

We will run your code multiple times till the time limit is met, and take the best solution amongst all the runs. You are provided with 10 sample problems. The final evaluation will be conducted on a similar set of problems. The points awarded will be your normalized performance score relative to other groups in the class. **To be fair to all, this part of the programming lab should be implemented in C.**

2. Multi-Agent Pacman[*]

This part of the lab has been adopted from Berkeley Pacman projects. In this project, you will design agents for the classic version of Pacman, including ghosts. Along the way, you will implement both minimax and expectimax search and try your hand at evaluation function design. The code base has not changed much from the previous lab, but please start with a fresh installation, rather than intermingling files from lab 1. As in lab 1, this lab includes an autograder for you to grade your answers on your machine. This can be run on all questions with the command:

python autograder.py

It can be run for one particular question, such as q2, by:

python autograder.py -q q2

It can be run for one particular test by commands of the form:

python autograder.py -t test_cases/q2/0-small-tree

By default, the *autograder* displays graphics with the -t option, but doesn't with the -q option. You can force graphics by using the --graphics flag, or force no graphics by using the --no-graphics flag.

The code for this project contains the following files, available as a zip archive. Files you'll edit:

multiAgents.py   Where all of your multi-agent search agents will reside.

---

[*] http://ai.berkeley.edu/multiagent.html

pacman.py   The main file that runs Pacman games. This file also describes a Pacman *GameState* type, which you will use extensively in this project

game.py   The logic behind how the Pacman world works. This file describes several supporting types like *AgentState, Agent, Direction,* and *Grid.*

util.py Useful data structures for implementing search algorithms.

Files you can ignore:

graphicsDisplay.py   Graphics for Pacman

graphicsUtils.py   Support for Pacman graphics

textDisplay.py ASCII graphics for Pacman

ghostAgents.py   Agents to control ghosts

keyboardAgents.py   Keyboard interfaces to control Pacman

layout.py   Code for reading layout files and storing their contents

autograder.py Project autograder

testParser.py Parses autograder test and solution files

testClasses.py General autograding test classes

test_cases/   Directory containing the test cases for each question

multiagentTestClasses.py   Lab 2 specific autograding test classes

You will fill in portions of multiAgents.py during the assignment. You should submit this file with your code and comments. Please do not change the other files in this distribution or submit any of our original files other than this file. Your code will be autograded for technical correctness. Please do not change the names of any provided functions or classes within the code, or you will wreak havoc on the *autograder.* However, the correctness of your implementation -- not the autograder's judgements -- will be the final judge of your score. If necessary, we will review and grade assignments individually to ensure that you receive due credit for your work.

Multi-Agent Pacman

First, play a game of classic Pacman:

python pacman.py

Now, run the provided *ReflexAgent* in *multiAgents.py*:

python pacman.py -p ReflexAgent

Note that it plays quite poorly even on simple layouts:

python pacman.py -p ReflexAgent -l testClassic

Inspect its code (in multiAgents.py) and make sure you understand what it's doing.

Question 1 (4 points): Reflex Agent

Improve the *ReflexAgent* in *multiAgents.py* to play respectably. The provided reflex agent code provides some helpful examples of methods that query the *GameState* for information. A capable reflex agent must consider both food locations and ghost locations to perform well. Your agent should easily and reliably clear the *testClassic* layout:

python pacman.py -p ReflexAgent -l testClassic

Try out your reflex agent on the default *mediumClassic* layout with one ghost or two (and animation off to speed up the display):

python pacman.py --frameTime 0 -p ReflexAgent -k 1
python pacman.py --frameTime 0 -p ReflexAgent -k 2

How does your agent fare? It will likely often die with 2 ghosts on the default board, unless your evaluation function is quite good.

Note: As features, try the reciprocal of important values (such as distance to food) rather than just the values themselves.

Note: The evaluation function you're writing is evaluating state-action pairs; in later parts of the project, you'll be evaluating states.

Options: Default ghosts are random; you can also play for fun with slightly smarter directional ghosts using *-g DirectionalGhost*. If the randomness is preventing you from telling whether your agent is improving, you can use -f to run with a fixed random seed (same random choices every game). You can also play multiple games in a row with *-n*. Turn off graphics with *-q* to run lots of games quickly.

Grading: we will run your agent on the *openClassic* layout 10 times. You will receive 0 points if your agent times out, or never wins. You will receive 1 point if your agent wins at least 5 times, or 2 points if your agent wins all 10 games. You will receive an addition 1 point if your agent's average score is greater than 500, or 2 points if it is greater than 1000. You can try your agent out under these conditions with

python autograder.py -q q1

To run it without graphics, use:
python autograder.py -q q1 --no-graphics

Question 2 (5 points): Minimax

Now you will write an adversarial search agent in the provided *MinimaxAgent* class stub in *multiAgents.py*. Your minimax agent should work with any number of ghosts, so you'll have to write an algorithm that is slightly more general than what you've previously seen in lecture. Your minimax tree will have multiple min layers (one for each ghost) for every max layer.

Your code should also expand the game tree to an arbitrary depth. Score the leaves of your minimax tree with the supplied *self.evaluationFunction*, which defaults to *scoreEvaluationFunction*. *MinimaxAgent* extends *MultiAgentSearchAgent*, which gives access to *self.depth* and *self.evaluationFunction*. Make sure your minimax code makes reference to these two variables where appropriate as these variables are populated in response to command line options.

Important: A single search ply is one Pacman move and all the ghosts' responses, so depth 2 search will involve Pacman and each ghost moving two times.

We will be checking your code to determine whether it explores the correct number of game states. This is the only reliable way to detect some very subtle bugs in implementations of minimax. As a result, the autograder will be very picky about how many times you call *GameState.generateSuccessor*. If you call it any more or less than necessary, the autograder will complain.

To test and debug your code, run

python autograder.py -q q2

This will show what your algorithm does on a number of small trees, as well as a pacman game. To run it without graphics, use:

python autograder.py -q q2 --no-graphics

Hints and Observations
The correct implementation of minimax will lead to Pacman losing the game in some tests. This is not a problem: as it is correct behaviour, it will pass the tests. The evaluation function for the pacman test in this part is already written (*self.evaluationFunction*). You shouldn't change this function, but recognize that now we're evaluating *states* rather than actions, as we were for the reflex agent. Look-ahead agents evaluate future states whereas reflex agents evaluate actions from the current state. The minimax values of the initial state in the *minimaxClassic* layout are 9,

8, 7, -492 for depths 1, 2, 3 and 4 respectively. Note that your minimax agent will often win (665/1000 games for us) despite the dire prediction of depth 4 minimax.

python pacman.py -p MinimaxAgent -l minimaxClassic -a depth=4

Pacman is always agent 0, and the agents move in order of increasing agent index. All states in minimax should be *GameStates*, either passed in to *getAction* or generated via *GameState.generateSuccessor*. In this project, you will not be abstracting to simplified states.
On larger boards such as *openClassic* and *mediumClassic* (the default), you'll find Pacman to be good at not dying, but quite bad at winning. He'll often thrash around without making progress. He might even thrash around right next to a dot without eating it because he doesn't know where he'd go after eating that dot. Don't worry if you see this behavior, question 5 will clean up all of these issues.

When Pacman believes that his death is unavoidable, he will try to end the game as soon as possible because of the constant penalty for living. Sometimes, this is the wrong thing to do with random ghosts, but minimax agents always assume the worst:

python pacman.py -p MinimaxAgent -l trappedClassic -a depth=3

Make sure you understand why Pacman rushes the closest ghost in this case.

Question 3 (5 points): Alpha-Beta Pruning

Make a new agent that uses alpha-beta pruning to more efficiently explore the minimax tree, in *AlphaBetaAgent.* Again, your algorithm will be slightly more general than the pseudocode from lecture, so part of the challenge is to extend the alpha-beta pruning logic appropriately to multiple minimizer agents.

You should see a speed-up (perhaps depth 3 alpha-beta will run as fast as depth 2 minimax). Ideally, depth 3 on *smallClassic* should run in just a few seconds per move or faster.

python pacman.py -p AlphaBetaAgent -a depth=3 -l smallClassic

The *AlphaBetaAgent* minimax values should be identical to the *MinimaxAgent* minimax values, although the actions it selects can vary because of different tie-breaking behavior. Again, the minimax values of the initial state in the *minimaxClassic* layout are 9, 8, 7 and -492 for depths 1, 2, 3 and 4 respectively.

Because we check your code to determine whether it explores the correct number of states, it is important that you perform alpha-beta pruning without reordering children. In other words, successor states should always be processed in the order returned by

GameState.getLegalActions. Again, do not call *GameState.generateSuccessor* more than necessary.

You must not prune on equality in order to match the set of states explored by our *autograder.* (Indeed, alternatively, but incompatible with our *autograder,* would be to also allow for pruning on equality and invoke alpha-beta once on each child of the root node, but this will not match the *autograder.*)

To test and debug your code, run

python autograder.py -q q3

This will show what your algorithm does on a number of small trees, as well as a pacman game. To run it without graphics, use:

python autograder.py -q q3 --no-graphics

The correct implementation of alpha-beta pruning will lead to Pacman losing some of the tests. This is not a problem: as it is correct behaviour, it will pass the tests.

Question 4 (6 points): Expectimax

Minimax and alpha-beta are great, but they both assume that you are playing against an adversary who makes optimal decisions. As anyone who has ever won tic-tac-toe can tell you, this is not always the case. In this question you will implement the *ExpectimaxAgent*, which is useful for modeling probabilistic behavior of agents who may make suboptimal choices.

As with the search and constraint satisfaction problems covered so far in this class, the beauty of these algorithms is their general applicability. To expedite your own development, we've supplied some test cases based on generic trees. You can debug your implementation on small the game trees using the command:

python autograder.py -q q4

Debugging on these small and manageable test cases is recommended and will help you to find bugs quickly. Make sure when you compute your averages that you use floats. Integer division in Python truncates, so that 1/2 = 0, unlike the case with floats where 1.0/2.0 = 0.5.

Once your algorithm is working on small trees, you can observe its success in Pacman. Random ghosts are of course not optimal minimax agents, and so modeling them with minimax search may not be appropriate. *ExpectimaxAgent,* will no longer take the min over all ghost actions, but the expectation according to your agent's model of how the ghosts act. To simplify your code,

assume you will only be running against an adversary which chooses amongst their getLegalActions uniformly at random.

To see how the *ExpectimaxAgent* behaves in Pacman, run:

python pacman.py -p ExpectimaxAgent -l minimaxClassic -a depth=3

You should now observe a more cavalier approach in close quarters with ghosts. In particular, if Pacman perceives that he could be trapped but might escape to grab a few more pieces of food, he'll at least try. Investigate the results of these two scenarios:

python pacman.py -p AlphaBetaAgent -l trappedClassic -a depth=3 -q -n 10
python pacman.py -p ExpectimaxAgent -l trappedClassic -a depth=3 -q -n 10

You should find that your *ExpectimaxAgent* wins about half the time, while your *AlphaBetaAgent* always loses. Make sure you understand why the behavior here differs from the minimax case.

The correct implementation of expectimax will lead to Pacman losing some of the tests. This is not a problem: as it is correct behaviour, it will pass the tests.

Question 5 (5 points): Evaluation Function

Write a better evaluation function for pacman in the provided function *betterEvaluationFunction.* The evaluation function should evaluate states, rather than actions like your reflex agent evaluation function did. You may use any tools at your disposal for evaluation, including your search code from the last project. With depth 2 search, your evaluation function should clear the *smallClassic* layout with one random ghost more than half the time and still run at a reasonable rate (to get full credit, Pacman should be averaging around 1000 points when he's winning).

python autograder.py -q q5

Grading: the *autograder* will run your agent on the *smallClassic* layout 10 times. We will assign points to your evaluation function in the following way:

If you win at least once without timing out the *autograder,* you receive 1 points. Any agent not satisfying these criteria will receive 0 points.
+1 for winning at least 5 times, +2 for winning all 10 times
+1 for an average score of at least 500, +2 for an average score of at least 1000 (including scores on lost games)
+1 if your games take on average less than 30 seconds on the *autograder* machine.

The additional points for average score and computation time will only be awarded if you win at least 5 times.

Hints and Observations

As for your reflex agent evaluation function, you may want to use the reciprocal of important values (such as distance to food) rather than the values themselves. One way you might want to write your evaluation function is to use a linear combination of features. That is, compute values for features about the state that you think are important, and then combine those features by multiplying them by different values and adding the results together. You might decide what to multiply each feature by based on how important you think it is.